

# Data Wrangling & Species Richness

*Jeff Oliver*

*31 October, 2019*

[INTRODUCTORY SENTENCE]

## Learning objectives

1. Demonstrate how to organize data in ‘tidy’ format
2. Develop quality control processes for data
3. Visualize measures of species richness and species diversity

[DESCRIPTION OR MOTIVATION; 2-4 sentences that would be used for an announcement]

---

## Getting started

Happy families are all alike; every unhappy family is unhappy in its own way - Leo Tolstoy

## Tidy data

After Wickham 2014

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

Open the file in Excel. What can be done to make these data “tidy”?

First, the data are effectively broken in two, with the most recent observations separate from observations from preceding days. So in some cases, a row has more than one observation. To fix this, we move the cells from the preceding observations to rows after the most recent observations.

We also want to:

- Delete any columns without data
- Delete any rows without data
- Add a column title for date
- Make sure data for preceding observations lines up with column title (currently the Count column has species names and the Species column has count data)

Let’s save this file as a CSV file so we can read it into R. We don’t want to overwrite the original data file, so we’ll call our file sweetwater-data-clean.csv.

---

## Data wrangling

Now that the data are in CSV format, we can read it into R and take a look. We want to keep a record of all our work, so we are going to type commands into a script file and save the script. As we did before, we add a few lines of comments at the top of our script with some relevant information.

```
# Analyze sweetwater bird species richness
# Jeff Oliver
# jcoliver@email.arizona.edu
# 2019-11-01

bird_data <- read.csv(file = "data/sweetwater-data-clean.csv")
```

We can look at the first few rows with the `head` function:

```
head(bird_data)
```

##	Date	Site	Count	Species
## 1	October 4th	1	1	Mallard Duck
## 2		NA	1	Red-winged blackbird
## 3		NA	1	Gila woodpecker
## 4		NA	1	Finch
## 5		NA	1	Mourning Dove
## 6		NA	1	Phainopepla

And we see that in rows two through 6 there aren't any values for the Date and Site columns. This is a prime example of the difference between how a human interprets data and how a computer interprets the data. Let's open up our CSV file in Excel again. As humans, we can infer that the value for "Date" should be October 4th all the way to row 35. But the computer is stupid. The computer looks at row 3, sees no date, and infers that the date is *missing*. We need to explicitly tell the computer what the value for date is. Before we do that, we have the opportunity to go ahead and also change the date to the standard ISO date format of "YYYY-MM-DD". In this case, October 4th becomes 2019-10-04.

Go ahead and update all the dates and fill in values appropriately. You should also note that for observations prior to October 10th, there is an offset that (1) needs to be fixed and (2) prompts deletion of now empty rows.

Also, for the 2019-10-04 observations, do the same thing with data in the Site column - that is, fill in values where appropriate. Save the file as a CSV (you can use the same filename, `sweetwater-data-clean.csv`) and read it into R.

```
# Analyze sweetwater bird species richness
# Jeff Oliver
# jcoliver@email.arizona.edu
# 2019-11-01

bird_data <- read.csv(file = "data/sweetwater-data-clean.csv")
```

Looking at the updated data, we see those first six rows no longer have missing data.

```
head(bird_data)
```

##	Date	Site	Count	Species
## 1	2019-10-04	1	1	Mallard Duck
## 2	2019-10-04	1	1	Red-winged blackbird
## 3	2019-10-04	1	1	Gila woodpecker
## 4	2019-10-04	1	1	Finch
## 5	2019-10-04	1	1	Mourning Dove
## 6	2019-10-04	1	1	Phainopepla

We can also use the `summary` function to get an idea of the data set overall:

```
summary(bird_data)
```

##	Date	Site	Count	Species
##	2019-09-28:20	Min. :1.000	Min. : 1.000	Gila Woodpecker : 3
##	2019-09-29:29	1st Qu.:2.000	1st Qu.: 1.000	Mallard Duck : 3
##	2019-09-30:34	Median :3.000	Median : 1.000	mourning dove : 3
##	2019-10-01: 5	Mean :2.941	Mean : 4.164	Blue Winged Teal: 2
##	2019-10-02:11	3rd Qu.:4.000	3rd Qu.: 3.000	Cooper's Hawk : 2
##	2019-10-03: 1	Max. :6.000	Max. :80.000	Gambils quail : 2
##	2019-10-04:34	NA's :100		(Other) :119

There are still missing values in the Site column, but that's OK - those are observations from preceding days, which don't have any site data.

---

## Species richness

OK, so now let's actually look at species richness. We are going to look at each date separately, and count the total number of species observed. To do this, we are going to use an additional package called `dplyr`; we will need to install the package and load it into memory before we try to use it. First, we check to see if the package is installed by looking at the Packages tab in the lower-right window. If `dplyr` is installed, we don't need to install it again. If `dplyr` is not listed, you can install it through the R console with the command:

```
install.packages("dplyr")
```

Now that the package is installed, we also have to tell R to load the package into memory, so we add a call to `library` to our script. We often add all `library` commands to the top of our script.

```
# Analyze sweetwater bird species richness
# Jeff Oliver
# jcoliver@email.arizona.edu
# 2019-11-01

library(dplyr)
bird_data <- read.csv(file = "data/sweetwater-data-clean.csv")
```

You might see some read messages when you load `dplyr`. As long as you don't see the word "Error" everything should be fine.

Now we can use some `dplyr` functions to calculate the total number of species seen on each day. We use the `group_by` function to first create a different pile of data for each day, then count the number of rows using the `n` function.

```
# Calculate species richness per day
richness <- bird_data %>%
  group_by(Date) %>%
  summarize(Total = n())
```

Let's look at the output by typing the name of our variable into the console:

```
richness
```

```
## # A tibble: 7 x 2
##   Date      Total
##   <fct>    <int>
## 1 2019-09-28    20
```

```
## 2 2019-09-29    29
## 3 2019-09-30    34
## 4 2019-10-01     5
## 5 2019-10-02    11
## 6 2019-10-03     1
## 7 2019-10-04    34
```

Pretty cool. We can even plot species richness for each day, too. We are going to use the ggplot2 package again, so first check to see if that package is installed. If not, install it by running `install.packages("ggplot2")` in the console, then update the script so that we have a call to `library(ggplot2)` at the top. Note we have to actually *run* this line in order to load ggplot2 into memory; just adding the code to script alone won't load it up.

```
# Analyze sweetwater bird species richness
# Jeff Oliver
# jcoliver@email.arizona.edu
# 2019-11-01


library(dplyr)
library(ggplot2)
bird_data <- read.csv(file = "data/sweetwater-data-clean.csv")
```

Again, if you see red messages when loading ggplot2, that's OK, so long as the word "Error" doesn't show up.

To plot richness for each day,

```
# Plot daily species richness
richness_plot <- ggplot(data = richness,
                        mapping = aes(x = Date,
                                      y = Total)) +

  geom_point()
print(richness_plot)
```



Hey, a plot!

Let's take a moment though to think about the data that went into making these plots. In fact, let's look at the first six rows again:

```
head(bird_data)
```

```
##      Date Site Count      Species
## 1 2019-10-04    1     1    Mallard Duck
## 2 2019-10-04    1     1 Red-winged blackbird
## 3 2019-10-04    1     1    Gila woodpecker
## 4 2019-10-04    1     1         Finch
## 5 2019-10-04    1     1    Mourning Dove
## 6 2019-10-04    1     1    Phainopepla
```

Hmmm...row four has "Finch", which isn't a species. At Sweetwater Wetlands, House Finches and Lesser Goldfinches are common, and sometimes other rare finches show up. Another example is on row 17. We can take a look at this part of the data by specifying which rows to show. Here we ask for R to print out rows 15 through 20:

```
bird_data[15:20, ]
```

```
##      Date Site Count      Species
## 15 2019-10-04    3     1    Gambils quail
## 16 2019-10-04    3     1    Phainopepla
## 17 2019-10-04    3     2    Finch sp
## 18 2019-10-04    3     1    Northern harrier
## 19 2019-10-04    3     1    Unidentified raptor
## 20 2019-10-04    3     3 Humming Bird: Rubythroated, Black chin
```

In addition to “Finch sp” we all so see two more examples of observations not identified to species. Because we don’t actually know what species these were, we will need to exclude them from our analyses.

---

## Computers are stupid

### Or, why quality control matters

We do not actually want to delete these records from our data, but rather we need a means of indicating whether or not they are identified to the species level. This should probably be done in the `sweetwater-data-clean.csv` file we’ve been working with. So open that file in Excel and add another column, immediately to the right of Species. Call this one “ID\_to\_species” and fill in values of 1 for those identified to species (e.g. Gila woodpecker would be scored as a 1) and values of 0 for those not identified to species (e.g. Finch sp and duck sp. would be scored as a 0). Save the file with the same name as before (`sweetwater-data-clean.csv`).

```
# Analyze sweetwater bird species richness
# Jeff Oliver
# jcoliver@email.arizona.edu
# 2019-11-01

library(dplyr)
library(ggplot2)
bird_data <- read.csv(file = "data/sweetwater-data-clean.csv")
```

Now we can update our richness calculation to only include observations that were identified to species. We do this by adding a filter to the process. Using the code we had before, before calling `group_by`, we use the `filter` function to only include those observations that have a 1 in the `ID_to_species` column.

```
# Calculate species richness per day
richness <- bird_data %>%
  filter(ID_to_species == 1) %>%
  group_by(Date) %>%
  summarize(Total = n())
```

We can compare the old calculations for richness

```
## # A tibble: 7 x 2
##   Date      Total
##   <fct>     <int>
## 1 2019-09-28     20
## 2 2019-09-29     29
## 3 2019-09-30     34
## 4 2019-10-01      5
## 5 2019-10-02     11
## 6 2019-10-03      1
## 7 2019-10-04     34
```

to the new ones:


```
richness
```

```
## # A tibble: 7 x 2
##   Date      Total
##   <fct>     <int>
## 1 2019-09-28     17
## 2 2019-09-29     26
```

```
## 3 2019-09-30    31
## 4 2019-10-01     4
## 5 2019-10-02     8
## 6 2019-10-03     1
## 7 2019-10-04    27
```

And note the numbers have dropped a little when we exclude those observations. Now when we run the code for plotting, we see the results changed, too:

```
richness_plot <- ggplot(data = richness,
                        mapping = aes(x = Date,
                                      y = Total)) +
  geom_point()
print(richness_plot)
```



Let's take a quick look at the species that were observed on October 4, just to double check. We can do that in the console using the pipe (`%>%`) and the `filter` command:

```
bird_data %>% filter(Date == "2019-10-04")
```

##	Date	Site	Count	Species
## 1	2019-10-04	1	1	Mallard Duck
## 2	2019-10-04	1	1	Red-winged blackbird
## 3	2019-10-04	1	1	Gila woodpecker
## 4	2019-10-04	1	1	Finch
## 5	2019-10-04	1	1	Mourning Dove
## 6	2019-10-04	1	1	Phainopepla
## 7	2019-10-04	2	5	Song Sparrow

## 8	2019-10-04	2	2	Gila Woodpecker
## 9	2019-10-04	2	1	Mallard Duck
## 10	2019-10-04	2	1	Lucy's Warbler
## 11	2019-10-04	2	1	Vermillion Flycatcher
## 12	2019-10-04	2	3	Curve-bill Thrasher
## 13	2019-10-04	3	9	mourning dove
## 14	2019-10-04	3	1	Turkey vulture
## 15	2019-10-04	3	1	Gambils quail
## 16	2019-10-04	3	1	Phainopepla
## 17	2019-10-04	3	2	Finch sp
## 18	2019-10-04	3	1	Northern harrier
## 19	2019-10-04	3	1	Unidentified raptor
## 20	2019-10-04	3	3	Humming Bird: Rubythroated, Black chin
## 21	2019-10-04	3	1	Flycatcher vermilion
## 22	2019-10-04	3	2	Warbler
## 23	2019-10-04	4	29	mourning dove
## 24	2019-10-04	4	1	mallard
## 25	2019-10-04	4	2	Gila Woodpecker
## 26	2019-10-04	4	1	Coopers' hawk
## 27	2019-10-04	4	1	Redwing black bird
## 28	2019-10-04	4	1	Vireo
## 29	2019-10-04	4	3	house finch
## 30	2019-10-04	4	1	Redtail hawk
## 31	2019-10-04	4	2	Gambils quail
## 32	2019-10-04	5	10	mourning dove
## 33	2019-10-04	5	2	vulture
## 34	2019-10-04	6	1	white winged dove
##	ID_to_species			
## 1		1		
## 2		1		
## 3		1		
## 4		0		
## 5		1		
## 6		1		
## 7		1		
## 8		1		
## 9		1		
## 10		1		
## 11		1		
## 12		1		
## 13		1		
## 14		1		
## 15		1		
## 16		1		
## 17		0		
## 18		1		
## 19		0		
## 20		0		
## 21		1		
## 22		0		
## 23		1		
## 24		1		
## 25		1		
## 26		1		



```
## 27          1
## 28          0
## 29          1
## 30          1
## 31          1
## 32          1
## 33          0
## 34          1
```

We see that there are 34 rows. This lines up with our initial plot, that showed 34 species. But look a little closer, notably at rows 1 and 9. Both rows are observations of Mallard Ducks. So this species actually got counted (at least) twice. Because multiple sites were surveyed, We can't just count the number of rows in this data to get the number of species. We should actually do a little more data wrangling to get an accurate count of species for the October 4 sampling. What we are going to do is to create a new data object we'll call `bird_data_sums` that sums up the values for each species on each day.

```
# Sum species counts for each day
bird_data_sums <- bird_data %>%
  filter(ID_to_species == 1) %>%
  group_by(Date, Species) %>%
  summarise(Count = sum(Count))
```

Now we can again look at the observations for October 4:

```
bird_data_sums %>% filter(Date == "2019-10-04")
```

```
## # A tibble: 21 x 3
## # Groups:   Date [1]
##   Date      Species      Count
##   <fct>     <fct>      <int>
## 1 2019-10-04 Coopers' hawk      1
## 2 2019-10-04 Curve-bill Thrasher  3
## 3 2019-10-04 Flycatcher vermilion  1
## 4 2019-10-04 Gambils quail      3
## 5 2019-10-04 Gila woodpecker      1
## 6 2019-10-04 Gila Woodpecker      4
## 7 2019-10-04 house finch        3
## 8 2019-10-04 Lucy's Warbler      1
## 9 2019-10-04 mallard            1
## 10 2019-10-04 Mallard Duck       2
## # ... with 11 more rows
```


Now this doesn't show us all the data, but looking at the end of the output, we see that R is showing us 10 rows of data, with 13 more rows. So we have gone from 34 species, down to 23 species.

We can update our code to calculate richness and re-create our plot. Remember to replace `bird_data` with `bird_data_sums` in the first line of the code. We can also remove the `filter` function, because we already filtered out observations that were not identified to species when we created `bird_data_sums`.

```
# Calculate species richness per day
richness <- bird_data_sums %>%
  group_by(Date) %>%
  summarize(Total = n())

# Plot daily species richness
richness_plot <- ggplot(data = richness,
  mapping = aes(x = Date,
    y = Total)) +
```

```
geom_point()
print(richness_plot)
```



The only thing to really change in our plot is for October 4. OK, some of you might have noticed this already, but let's look again at those October 4 data:

```
bird_data_sums %>% filter(Date == "2019-10-04")
```

```
## # A tibble: 21 x 3
## # Groups:   Date [1]
##   Date      Species      Count
##   <fct>     <fct>     <int>
## 1 2019-10-04 Coopers' hawk      1
## 2 2019-10-04 Curve-bill Thrasher    3
## 3 2019-10-04 Flycatcher vermilion    1
## 4 2019-10-04 Gambils quail          3
## 5 2019-10-04 Gila woodpecker          1
## 6 2019-10-04 Gila Woodpecker          4
## 7 2019-10-04 house finch              3
## 8 2019-10-04 Lucy's Warbler            1
## 9 2019-10-04 mallard                  1
## 10 2019-10-04 Mallard Duck              2
## # ... with 11 more rows
```

So, there's another issue to address. Look at rows 9 and 10. We have a species that is represented twice. Why? Here is a great example of computers only being able to do what you tell them to. We told R to count the total number of individuals for each species. Unfortunately, R sees "mallard" and "Mallard Duck"

as two different species.

What do we do?

At this point, the best thing is probably to go into the Excel file and change the species names to be consistent. Now so far, we haven't actually *changed* any of the data - we added a column and we filled in values where appropriate, but we did not change anything. When changing data, you always want to be careful. At this point, we should be keeping track of changes we made to the data. We are going to do this in a separate file, with notes about what we did. So start by creating a new text file (File > New File > Text File) and adding:

```
README for sweetwater species data
2019-11-01
+ Created CSV from original data file
+ Changed to ISO date format
+ Standardized spelling of species' names
```

And save this file as "README.md".

Now we actually have to do the standardization of the species names, so open up the sweetwater-data-clean.csv file in Excel. What's the best way of going about changing names? We could search for non-standard spellings like "Gambils quail" but that assumes we know all the alternative spellings. What might help our cause is to sort by the Species column so things like "mallard" and "Mallard Duck" show up next to one another. **Note:** this sorting is going to rearrange the order of the rows, so we should *only* do this if the order of the rows does not matter. Remember the original data? Where the date was only shown in the first row? If we had sorted those data it would have been *bad*. Why? Because in that case, order *did* matter. But by using tidy data formats, we can re-order rows without losing any information for a single observation.

So go ahead and change names so they are consistent. This includes

- consistent spelling (Gambel's vs. Gambils),
- consistent capitalization (Gila Woodpecker vs. Gila woodpecker), and
- constant use of punctuation (Red-winged Blackbird vs. Redwing black bird and Cooper's Hawk vs. Coopers' Hawk).

Save the file to sweetwater-data-clean.csv and go back into R.

Re-run the line at the top of our script with the `read.csv` command:

```
bird_data <- read.csv(file = "data/sweetwater-data-clean.csv")
```

Re-run the line counting the number of each species per day:


```
# Sum species counts for each day
bird_data_sums <- bird_data %>%
  filter(ID_to_species == 1) %>%
  group_by(Date, Species) %>%
  summarise(Count = sum(Count))
```

And calculate richness for each day and draw the plot

```
# Calculate species richness per day
richness <- bird_data_sums %>%
  group_by(Date) %>%
  summarize(Total = n())

# Plot daily species richness
richness_plot <- ggplot(data = richness,
  mapping = aes(x = Date,
    y = Total)) +
```

```
geom_point()
print(richness_plot)
```



---

## Species richness

```
richness <- bird_data %>%
  group_by(Date) %>%
  summarize(Total = n())
```

---

## Species diversity

There's considerable data wrangling necessary here. `vegan::diversity` needs a wide-formatted data set, but since there were multiple observations of a single species on October 4, this creates challenges for converting long data to wide-formatted data (because date is effectively a key, but there are multiple observations of some species for 2019-10-04).

UPDATE: By summing counts for October 4, this should now be possible.

```
bird_data_wide <- pivot_wider(data = bird_data[, -2],
                             names_from = Species,
```

```

      values_from = Count,
      values_fill = list(Count = 0))

species_counts <- as.matrix(as.data.frame(bird_data_wide[, 3:119]))
diversity <- vegan::diversity(x = bird_data_wide[, 3:119])

```

Could alternatively do a species accumulation curve, using `duplicated`

```

first <- bird_data[!duplicated(bird_data$Species), ]
first <- first[order(first$Date), ]
first_date <- first %>%
  group_by(Date) %>%
  summarize(new_species = n())
first_date$cumulative <- cumsum(x = first_date$new_species)
first_date$Date <- as.Date(x = first_date$Date)
plot(x = first_date$Date,
     y = first_date$cumulative,
     type = "l")

```

Diversity

Shannon's index

$$H = - \sum_{i=1}^s p_i \ln p_i$$

is the sum over  $s$  species, where  $p_i$  is the proportion individuals of the  $i^{th}$  species ( $n_i$ ) out of the total number of individuals ( $N$ ).

```
use vegan::diversity(x = bird_data$Count, index = "Shannon")
```

---

## Additional resources

- Wickham, H. 2014. Tidy data. *The Journal of Statistical Software* **59** <http://www.jstatsoft.org/v59/i10/>
- A nice introduction to [calculating measures of species diversity](#)
- A [PDF version](#) of this lesson

---

Back to learn-r main page

Questions? e-mail me at [jcoliver@email.arizona.edu](mailto:jcoliver@email.arizona.edu).